



Addition generalization · feedback scale at a fixed learning-rate schedule


RLT-1 4+4




RLT-1 5+3





RLT-1 6+2



RLT-1 7+1



RLT-1 8+0



Seed 42 · 2,000 training steps
 Training: mixed 1-8 digits
 Same LR schedule: $1e-4 \rightarrow 5e-6$
 Global batch 512 · TBPTT128
 256 shared examples per length
 Best ID-loss checkpoint; earliest ties
 Single seed; no error bars