



Mod-5 without brackets · ID validation · final answer


Layer 4 + 4




Layer 5 + 3




Layer 6 + 2



Layer 7 + 1



Layer 8 + 0



- RLT-1
- RLT-0
- RLT-2 chunk4
- GPT Layer 8

5,000 training steps · global batch 512
All models: seeds 42, 43, 44
Mean ± sample SD; all seeds at each step
All runs completed training
SD clipped to 0-100% for display