



Parity · 16 layers · best versus final checkpoint

Best ID-loss checkpoint (earliest tie)



Fixed step-2,000 checkpoint



RLT-1 8+8 (1400) RLT-1 10+6 (1800) RLT-1 12+4 (1000) RLT-1 14+2 (1500) RLT-1 16+0 (1900)
 RLT-1 9+7 (1300) RLT-1 11+5 (1300) RLT-1 13+3 (2000) RLT-1 15+1 (1100) Transformer 16 (2000)

Seed 42 · all runs completed 2,000 steps · training batch 1,024 · 1,024 shared examples per length · best steps in parentheses