



Parity · 16 layers · best versus final checkpoint

Best ID-loss checkpoint (earliest tie)



Fixed step-2,000 checkpoint



- RLT-1 Layer 8 + 8 (1400)
- RLT-1 Layer 9 + 7 (1300)
- RLT-1 Layer 10 + 6 (1800)
- RLT-1 Layer 11 + 5 (1300)
- RLT-1 Layer 12 + 4 (1000)
- RLT-1 Layer 13 + 3 (2000)
- RLT-1 Layer 14 + 2 (1500)
- RLT-1 Layer 15 + 1 (1100)
- RLT-1 Layer 16 + 0 (1900)
- GPT Layer 16 (2000)

Seed 42 · all runs completed 2,000 steps · training batch 1,024 · 1,024 shared examples per length · best steps in parentheses