



Parity · length generalization · final answer


Layer 4 + 4




Layer 5 + 3




Layer 6 + 2



Layer 7 + 1



Layer 8 + 0



- RLT-1
- RLT-0
- RLT-2 chunk4
- GPT Layer 8

2,000 training steps · global batch 512
 All models: seeds 42, 43, 44
 Best ID-loss checkpoints; earliest ties
 Mean ± sample SD; full seed sets only
 1,024 shared examples per length
 SD clipped to 0-100% for display