






S5 swaps · length generalization · final state



- RLT-1
- RLT-0
- RLT-2 chunk4
- Transformer 8

2,000 training steps · global batch 512
 All models: seeds 42, 43, 44
 Best ID-loss checkpoints; earliest ties
 Mean ± sample SD; full seed sets only
 1,024 shared examples per length
 SD clipped to 0-100% for display